

2403A52031

AKULA MANOJ

B-02

Step 1: Install & Import Libraries

```
!pip install nltk spacy wordcloud

Requirement already satisfied: nltk in /usr/local/lib/python3.12/dist-
packages (3.9.1)
Requirement already satisfied: spacy in
/usr/local/lib/python3.12/dist-packages (3.8.11)
Requirement already satisfied: wordcloud in
/usr/local/lib/python3.12/dist-packages (1.9.5)
Requirement already satisfied: click in
/usr/local/lib/python3.12/dist-packages (from nltk) (8.3.1)
Requirement already satisfied: joblib in
/usr/local/lib/python3.12/dist-packages (from nltk) (1.5.3)
Requirement already satisfied: regex>=2021.8.3 in
/usr/local/lib/python3.12/dist-packages (from nltk) (2025.11.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (from nltk) (4.67.1)
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in
/usr/local/lib/python3.12/dist-packages (from spacy) (3.0.12)
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (1.0.5)
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (1.0.15)
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.0.13)
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in
/usr/local/lib/python3.12/dist-packages (from spacy) (3.0.12)
Requirement already satisfied: thinc<8.4.0,>=8.3.4 in
/usr/local/lib/python3.12/dist-packages (from spacy) (8.3.10)
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in
/usr/local/lib/python3.12/dist-packages (from spacy) (1.1.3)
Requirement already satisfied: srslly<3.0.0,>=2.4.3 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.5.2)
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.0.10)
Requirement already satisfied: weasel<0.5.0,>=0.4.2 in
/usr/local/lib/python3.12/dist-packages (from spacy) (0.4.3)
Requirement already satisfied: typer-slim<1.0.0,>=0.3.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (0.21.1)
Requirement already satisfied: numpy>=1.19.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.0.2)
Requirement already satisfied: requests<3.0.0,>=2.13.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.32.4)
```

```
Requirement already satisfied: pydantic!=1.8,!>=1.8.1,<3.0.0,>=1.7.4 in
/usr/local/lib/python3.12/dist-packages (from spacy) (2.12.3)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.12/dist-packages (from spacy) (3.1.6)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.12/dist-packages (from spacy) (75.2.0)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.12/dist-packages (from spacy) (25.0)
Requirement already satisfied: pillow in
/usr/local/lib/python3.12/dist-packages (from wordcloud) (11.3.0)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.12/dist-packages (from wordcloud) (3.10.0)
Requirement already satisfied: annotated-types>=0.6.0 in
/usr/local/lib/python3.12/dist-packages (from pydantic!=1.8,!>=1.8.1,<3.0.0,>=1.7.4->spacy) (0.7.0)
Requirement already satisfied: pydantic-core==2.41.4 in
/usr/local/lib/python3.12/dist-packages (from pydantic!=1.8,!>=1.8.1,<3.0.0,>=1.7.4->spacy) (2.41.4)
Requirement already satisfied: typing-extensions>=4.14.1 in
/usr/local/lib/python3.12/dist-packages (from pydantic!=1.8,!>=1.8.1,<3.0.0,>=1.7.4->spacy) (4.15.0)
Requirement already satisfied: typing-inspection>=0.4.2 in
/usr/local/lib/python3.12/dist-packages (from pydantic!=1.8,!>=1.8.1,<3.0.0,>=1.7.4->spacy) (0.4.2)
Requirement already satisfied: charset_normalizer<4,>=2 in
/usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.13.0->spacy) (3.4.4)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.13.0->spacy) (3.11)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.13.0->spacy) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.13.0->spacy) (2026.1.4)
Requirement already satisfied: blis<1.4.0,>=1.3.0 in
/usr/local/lib/python3.12/dist-packages (from thinc<8.4.0,>=8.3.4->spacy) (1.3.3)
Requirement already satisfied: confection<1.0.0,>=0.0.1 in
/usr/local/lib/python3.12/dist-packages (from thinc<8.4.0,>=8.3.4->spacy) (0.1.5)
Requirement already satisfied: cloudpathlib<1.0.0,>=0.7.0 in
/usr/local/lib/python3.12/dist-packages (from weasel<0.5.0,>=0.4.2->spacy) (0.23.0)
Requirement already satisfied: smart-open<8.0.0,>=5.2.1 in
/usr/local/lib/python3.12/dist-packages (from weasel<0.5.0,>=0.4.2->spacy) (7.5.0)
Requirement already satisfied: MarkupSafe>=2.0 in
```

```

/usr/local/lib/python3.12/dist-packages (from jinja2->spacy) (3.0.3)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(1.3.3)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(4.61.1)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(1.4.9)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(3.3.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->wordcloud)
(2.9.0.post0)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7-
>matplotlib->wordcloud) (1.17.0)
Requirement already satisfied: wrapt in
/usr/local/lib/python3.12/dist-packages (from smart-
open<8.0.0,>=5.2.1->weasel<0.5.0,>=0.4.2->spacy) (2.0.1)

import pandas as pd
import re
import nltk
import matplotlib.pyplot as plt

from nltk.corpus import stopwords
from sklearn.feature_extraction.text import TfidfVectorizer
from wordcloud import WordCloud

nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.

True

```

Step 2: Load Dataset

```

df = pd.read_csv("/content/Tweets.csv")
df = df[['text', 'airline_sentiment']]
df.head()

{"summary": "{\n    \"name\": \"df\", \n    \"rows\": 14640, \n    \"fields\": [\n        {\n            \"column\": \"text\", \n            \"properties\": {\n
```

```

\"dtype\": \"string\", \n      \"num_unique_values\": 14427, \n
\"samples\": [ \n          \"@JetBlue so technically I could drive to \nJFK now and put in. Request for tomorrow's flight?\", \n
\"@united why I won't check my carry on. Watched a handler throw this \nbag -- miss the conveyer belt -- sat there 10 min \nhttp://t.co/lyoocx5mSH\", \n          \"@SouthwestAir you guys are so \nclever \\ud83d\\ude03 http://t.co/qn5odUGFqK\" \n      ], \n
\"semantic_type\": \"\", \n          \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"airline_sentiment\", \n
\"properties\": { \n      \"dtype\": \"category\", \n
\"num_unique_values\": 3, \n      \"samples\": [ \n        \"neutral\", \n          \"positive\", \n          \"negative\" \n      ], \n      \"semantic_type\": \"\", \n          \"description\": \"\" \n    } \n  } \n], \n  \"type\": \"dataframe\", \n  \"variable_name\": \"df\" \n}

```

Step 3: Text Preprocessing

```

stop_words = set(stopwords.words('english'))

def clean_tweet(text):
    text = text.lower()
    text = re.sub(r'http\S+', '', text)           # remove URLs
    text = re.sub(r'@\w+', '', text)              # remove mentions
    text = re.sub(r'#\w+', '', text)              # remove hashtags
    text = re.sub(r'[^a-z\s]', '', text)          # remove special chars
    tokens = text.split()
    tokens = [word for word in tokens if word not in stop_words]
    return " ".join(tokens)

df['clean_text'] = df['text'].apply(clean_tweet)
df.head()

{
  "summary": {
    "name": "df",
    "rows": 14640,
    "fields": [
      {
        "column": "text",
        "properties": {
          "dtype": "string",
          "num_unique_values": 14427,
          "samples": [
            "@JetBlue so technically I could drive to \nJFK now and put in. Request for tomorrow's flight?",
            "@united why I won't check my carry on. Watched a handler throw this \nbag -- miss the conveyer belt -- sat there 10 min \nhttp://t.co/lyoocx5mSH",
            "@SouthwestAir you guys are so \nclever \\ud83d\\ude03 http://t.co/qn5odUGFqK"
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "airline_sentiment",
        "properties": {
          "dtype": "category",
          "num_unique_values": 3,
          "samples": [
            "neutral",
            "positive",
            "negative"
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ],
    "type": "dataframe",
    "variable_name": "df"
}

```

```
\"num_unique_values\": 14046, \"samples\": [\n    \"rude grouchy agent dulles checkin got hour trip lousy start needs\n    nap something\", \"smart makes angry mt weve partnered\n    bring free digital access onboard\", \"really drunk someone\n    stop serving\"\n], \"semantic_type\": \"\", \"\n    \"description\": \"\\n        }\\n        }\\n        ]\\n    }\", \"type\": \"dataframe\", \"variable_name\": \"df\"}
```

Step 4: Filter Negative Tweets

```
negative_tweets = df[df['airline_sentiment'] == 'negative']\nnegative_tweets.shape\n(9178, 3)
```

Step 5: TF-IDF Vectorization

```
vectorizer = TfidfVectorizer(max_features=20)\ntfidf_matrix = vectorizer.fit_transform(negative_tweets['clean_text'])\n\ntfidf_df = pd.DataFrame(\n    tfidf_matrix.toarray(),\n    columns=vectorizer.get_feature_names_out()\n)\n\ntfidf_df.mean().sort_values(ascending=False)\n\nflight      0.161364\nget         0.068515\ncancelled   0.053937\nservice     0.052331\ntime        0.044687\nhours       0.044553\nhelp        0.043687\nim          0.042881\ncustomer    0.042700\nhold        0.040962\nplane       0.039804\nus          0.038649\nstill       0.037706\ndelayed     0.037145\ncant        0.036345\none         0.034583\namp          0.034461\nhour        0.033592\ncall        0.032677\nflightled   0.027828\ndtype: float64
```

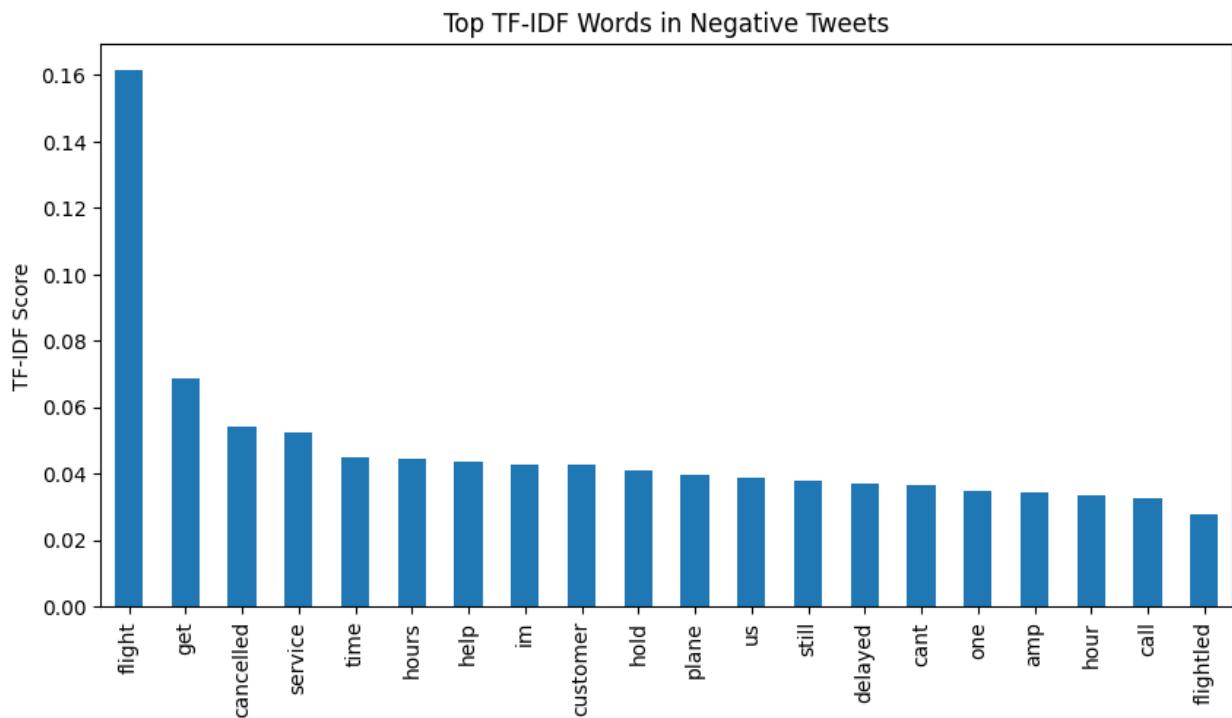
Step 6: Bar Chart (Top Negative Words)

```

top_words = tfidf_df.mean().sort_values(ascending=False)

plt.figure(figsize=(10,5))
top_words.plot(kind='bar')
plt.title("Top TF-IDF Words in Negative Tweets")
plt.ylabel("TF-IDF Score")
plt.show()

```



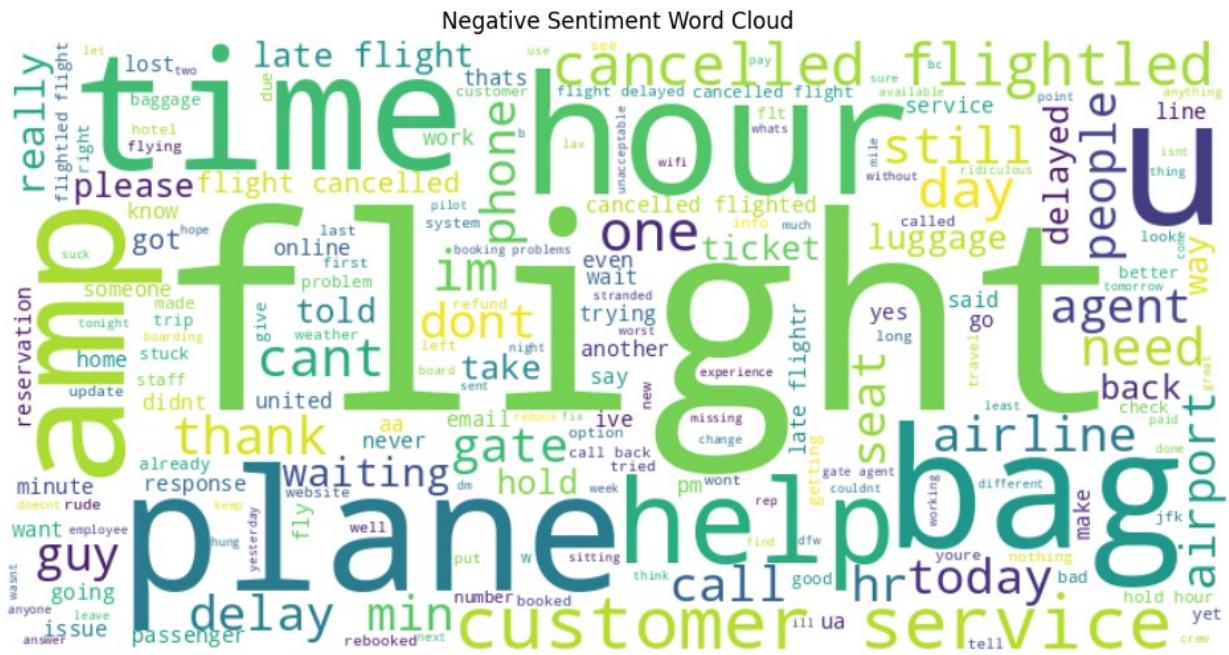
Step 7: Word Cloud

```

wordcloud = WordCloud(
    width=800,
    height=400,
    background_color='white'
).generate(" ".join(negative_tweets['clean_text']))

plt.figure(figsize=(12,6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title("Negative Sentiment Word Cloud")
plt.show()

```



Discussion

This assignment applied NLP preprocessing techniques to clean and prepare Twitter data for analysis. TF-IDF was used to identify important words contributing to negative airline sentiment. The TF-IDF matrix enabled quantitative comparison of terms across tweets. Bar chart and word cloud visualizations clearly highlighted frequent negative sentiment keywords. The results demonstrate the effectiveness of TF-IDF in extracting meaningful insights from social media text.